

# Case Study - AI Legal & Compliance platform

## Your task

Build a workflow assistant for a Legal and Compliance team that automates the contract management and regulatory monitoring processes. The assistant should ingest documents, extract key information, update internal systems, monitor regulatory changes and trigger downstream actions.

### Automated contract intake and metadata extraction

- Build a website where people can upload contracts. Once uploaded, extract metadata from the PDF in a standardized format. The metadata includes: parties, contract type, notice period, renewal date, costs, etc.
  - As soon as a new contract is uploaded, the compliance manager should also receive a Slack message to review it. The stage here is “Under Review”
  - The compliance manager clicks on the link on Slack to review the contract, which they can mark “Reviewed” on the website.
- Build an Airtable that will store this information and also show it in the website in an Airtable so that the compliance team can have visibility into existing contracts.
  - There will be two stages in the contract table: “Reviewed” and “Under Review”.
- Compute renewal/termination deadlines from the extracted metadata and schedule Slack reminders for responsible owners. Do this 1 month before and 1 week before the deadline.

### Compliance Dashboard

- Aggregate information on contract counts, upcoming renewals, policy updates and compliance findings into a dashboard. Provide filters by counterparty, cost, jurisdiction and risk.
- A Q&A chatbot where the compliance team can ask questions about the underlying data, like “How many contract renewals are upcoming in the next 2 weeks? Which ones specifically?” The questions can be based on both structured metadata in the table and unstructured contents of documents. The sources must be clearly cited in both cases.

### Regulatory Monitoring

- At scheduled intervals (e.g. daily), use web-search and retrieval to scan official DE/EU/US regulatory feeds for changes relevant to Bit Capital’s operations.
  - Regulatory updates include pieces of legislation and publications from regulators related to: DORA, AML, AI, data protection, capital markets, investment firms, ESG, crypto, etc.

- Pull the new documents and store them in a vector store for long-term querying. One should be able to query these documents on the compliance website via the Q&A chatbot, with the sources clearly cited.
- Send a Slack message in a compliance channel on material new information.
- A weekly summary of the new regulatory updates that can be read on the compliance website and be exported as a PDF.

## Things we will watch // Deliverables

- Source code with clear modular design and “good primitives” (functions for ingestion, vector-store operations, agent orchestration and API calls).
  - Optional short Loom/video demo.
- An explanation of the architecture, the technology stack chosen and why, including model choice in different use cases.
- Prompts used for metadata extraction and policy updates, with explanations of design choices and fallback strategies.
- A final report that summarizes performance metrics and lessons learned. A critical component is the analysis of model evaluation and tokenomics. Specifically, the report should approximate token usage for each workflow and explain how this factor influenced model selection, including cost-quality trade-offs.